

SDAIA BOOTCAMP PROJECT - Google Play Store

January 12, 2022

1 1. Introduction

With technology dominating the modern world; the app industry has become a high thriving. New apps are coming to the market every while. Indeed, app usage is still growing at a steady rate, but some are higher than others. There are different demands for different apps based on several features. Data science potentials can be utilized to drive app-making businesses and app developers to the right road.

2 2. Dataset Description

Google Play Store is a big digital distribution service that provides apps supported by Android-certified devices and Chrome OS. We found a dataset on Kaggle that contains data of 10k Play Store apps for analyzing the Android market. The dataset has 13 columns -which will be shown in the following subsection- and 10842 rows. It is aiming to use these apps' statistics to predict which apps are more likely to be installed or get a high rate.

2.1 2.1. Sneak Peek on the Dataset

This is the link to the dataset on Kaggle (<https://www.kaggle.com/lava18/google-play-store-apps>)

2.1.1 2.1.1. Importing Libraries and the Dataset

```
[1]: #import the libraries
import numpy as np
import pandas as pd

#read the dataset
df = pd.read_csv('googleplaystore.csv')
df.head(10)
```

```
[1]:
```

	App	Category	Rating \
0	Photo Editor & Candy Camera & Grid & ScrapBook	ART_AND_DESIGN	4.1
1	Coloring book moana	ART_AND_DESIGN	3.9
2	U Launcher Lite - FREE Live Cool Themes, Hide ...	ART_AND_DESIGN	4.7
3	Sketch - Draw & Paint	ART_AND_DESIGN	4.5
4	Pixel Draw - Number Art Coloring Book	ART_AND_DESIGN	4.3
5	Paper flowers instructions	ART_AND_DESIGN	4.4
6	Smoke Effect Photo Maker - Smoke Editor	ART_AND_DESIGN	3.8

7	Infinite Painter	ART_AND_DESIGN	4.1
8	Garden Coloring Book	ART_AND_DESIGN	4.4
9	Kids Paint Free - Drawing Fun	ART_AND_DESIGN	4.7

	Reviews	Size	Installs	Type	Price	Content	Rating \
0	159	19M	10,000+	Free	0	Everyone	
1	967	14M	500,000+	Free	0	Everyone	
2	87510	8.7M	5,000,000+	Free	0	Everyone	
3	215644	25M	50,000,000+	Free	0	Teen	
4	967	2.8M	100,000+	Free	0	Everyone	
5	167	5.6M	50,000+	Free	0	Everyone	
6	178	19M	50,000+	Free	0	Everyone	
7	36815	29M	1,000,000+	Free	0	Everyone	
8	13791	33M	1,000,000+	Free	0	Everyone	
9	121	3.1M	10,000+	Free	0	Everyone	

	Genres	Last Updated	Current Ver \
0	Art & Design	January 7, 2018	1.0.0
1	Art & Design;Pretend Play	January 15, 2018	2.0.0
2	Art & Design	August 1, 2018	1.2.4
3	Art & Design	June 8, 2018	Varies with device
4	Art & Design;Creativity	June 20, 2018	1.1
5	Art & Design	March 26, 2017	1.0
6	Art & Design	April 26, 2018	1.1
7	Art & Design	June 14, 2018	6.1.61.1
8	Art & Design	September 20, 2017	2.9.2
9	Art & Design;Creativity	July 3, 2018	2.8

	Android Ver
0	4.0.3 and up
1	4.0.3 and up
2	4.0.3 and up
3	4.2 and up
4	4.4 and up
5	2.3 and up
6	4.0.3 and up
7	4.2 and up
8	3.0 and up
9	4.0.3 and up

```
[2]: df.shape
```

```
[2]: (10841, 13)
```

2.2. Features Explanation

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10841 entries, 0 to 10840
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   App                   10841 non-null  object
1   Category              10841 non-null  object
2   Rating                9367 non-null   float64
3   Reviews               10841 non-null  object
4   Size                  10841 non-null  object
5   Installs              10841 non-null  object
6   Type                  10840 non-null  object
7   Price                 10841 non-null  object
8   Content Rating        10840 non-null  object
9   Genres                10841 non-null  object
10  Last Updated          10841 non-null  object
11  Current Ver           10833 non-null  object
12  Android Ver           10838 non-null  object
dtypes: float64(1), object(12)
memory usage: 1.1+ MB
```

1. App: is for the name of the app.
2. Category: means the category of a certain app
3. Rating: displays the rating of the app in Google Play Store on 5 Point numerical rating scale.
4. Reviews: this shows the number of reviews given to the app.
5. Size: shows the size of the app in megabytes.
6. Installs: shows the number of times the app got installed
7. Type: this shows if the app is free or paid.
8. Price: shows the price of the app in dollars. If the app is free the value will be 0.
9. Content Rating: shows the rating of the content if it's for everyone or specified for a specific audience.
10. Genres: means the genre of a certain app (it appears to be the same as the Category feature)
11. Last Updated: shows the date of the last update of the app.
12. Current Ver: shows the number of the current version of the app.
13. Android Ver: shows the number of Android versions that the app support.

3. Purpose of the Project

There are two principle goals for this project. The first goal is to classify apps based on their Installs and Rating, taking into consideration other features like Category, Type, Price and Content Rating. We're aiming to run many experiments with different classifier models and many trails to discover their effect on the accuracy scores. The second goal of this project is to predict Rating based on set of app features like Category, Reviews, Size, Installs, Type, Price and Content Rating. We're aiming to build a regression model that can predict the app rating.

4 4. Data Cleaning

We plan to perform basic cleaning for the dataset and analyze some features. We will try to drop columns that will appear to be irrelevant to our analysis.

4.1 4.1. Removing Duplicates

```
[4]: df[df.duplicated(subset='App')]
```

```
[4]:
```

	App	Category	Rating	Reviews	\
229	Quick PDF Scanner + OCR FREE	BUSINESS	4.2	80805	
236	Box	BUSINESS	4.2	159872	
239	Google My Business	BUSINESS	4.4	70991	
256	ZOOM Cloud Meetings	BUSINESS	4.4	31614	
261	join.me - Simple Meetings	BUSINESS	4.0	6989	
...	
10715	FarmersOnly Dating	DATING	3.0	1145	
10720	Firefox Focus: The privacy browser	COMMUNICATION	4.4	36981	
10730	FP Notebook	MEDICAL	4.5	410	
10753	Slickdeals: Coupons & Shopping	SHOPPING	4.5	33599	
10768	AAFP	MEDICAL	3.8	63	

	Size	Installs	Type	Price	Content	Rating	\
229	Varies with device	5,000,000+	Free	0		Everyone	
236	Varies with device	10,000,000+	Free	0		Everyone	
239	Varies with device	5,000,000+	Free	0		Everyone	
256	37M	10,000,000+	Free	0		Everyone	
261	Varies with device	1,000,000+	Free	0		Everyone	
...	
10715	1.4M	100,000+	Free	0		Mature 17+	
10720	4.0M	1,000,000+	Free	0		Everyone	
10730	60M	50,000+	Free	0		Everyone	
10753	12M	1,000,000+	Free	0		Everyone	
10768	24M	10,000+	Free	0		Everyone	

	Genres	Last Updated	Current Ver	\
229	Business	February 26, 2018	Varies with device	
236	Business	July 31, 2018	Varies with device	
239	Business	July 24, 2018	2.19.0.204537701	
256	Business	July 20, 2018	4.1.28165.0716	
261	Business	July 16, 2018	4.3.0.508	
...	
10715	Dating	February 25, 2016	2.2	
10720	Communication	July 6, 2018	5.2	
10730	Medical	March 24, 2018	2.1.0.372	
10753	Shopping	July 30, 2018	3.9	
10768	Medical	June 22, 2018	2.3.1	

	Android Ver
229	4.0.3 and up
236	Varies with device
239	4.4 and up
256	4.0 and up
261	4.4 and up
...	...
10715	4.0 and up
10720	5.0 and up
10730	4.4 and up
10753	4.4 and up
10768	5.0 and up

[1181 rows x 13 columns]

There are 1181 duplicated apps apparently. They need to be dropped.

```
[5]: df.drop_duplicates(subset='App', inplace=True, ignore_index=True)
```

4.2 4.2. Dropping Irrelevant Columns

From looking at the head of the table, we can notice that “Category” and “Genres” features are a little similar.

```
[6]: df[['Category', 'Genres']]
```

```
[6]:
```

	Category	Genres
0	ART_AND_DESIGN	Art & Design
1	ART_AND_DESIGN	Art & Design;Pretend Play
2	ART_AND_DESIGN	Art & Design
3	ART_AND_DESIGN	Art & Design
4	ART_AND_DESIGN	Art & Design;Creativity
...
9655	FAMILY	Education
9656	FAMILY	Education
9657	MEDICAL	Medical
9658	BOOKS_AND_REFERENCE	Books & Reference
9659	LIFESTYLE	Lifestyle

[9660 rows x 2 columns]

```
[7]: print(df['Category'].nunique())
print(df['Genres'].nunique())
```

34
119

From the above, we can notice that “Category” is more abstract than the “Genres” feature. The “Genres” has more values as it takes into consideration more than one genre to the app. We want to be exclusive as possible so we will work on the “Category” column and ignore the “Genres” column.

For our classification and regression models, we consider “App”, “Last Updated”, “Current Ver”, and “Android Ver” in addition to “Genres” as irrelevant features. So, we need to drop all these irrelevant columns

```
[8]: df = df.drop(['App', 'Genres', 'Last Updated', 'Current Ver', 'Android Ver'],
    ↪axis = 1)
```

```
[9]: df.head(3)
```

```
[9]:
```

	Category	Rating	Reviews	Size	Installs	Type	Price	Content	Rating
0	ART_AND_DESIGN	4.1	159	19M	10,000+	Free	0		Everyone
1	ART_AND_DESIGN	3.9	967	14M	500,000+	Free	0		Everyone
2	ART_AND_DESIGN	4.7	87510	8.7M	5,000,000+	Free	0		Everyone

4.3 Handling Missing Values

```
[10]: #Detecting missing values
df.isna().sum()
```

```
[10]: Category          0
Rating          1463
Reviews          0
Size             0
Installs         0
Type             1
Price            0
Content Rating    1
dtype: int64
```

We can see that ‘Rating’ has a large number of rows with missing values. So, instead of dropping these rows, we will fill them with the mean value of the ‘Rating’.

```
[11]: #Filling missing values with mean
RatingMean = df['Rating'].mean()
df['Rating'] = df['Rating'].fillna(RatingMean)
```

As we have seen, ‘Type’ and ‘Content Rating’ have only one missing value for each, so, we will drop these rows.

```
[12]: #Dropping missing values
df = df.dropna(subset=['Type'])
df = df.dropna(subset=['Content Rating'])
```

```
[13]: df.isna().sum()
```

```
[13]: Category          0
      Rating            0
      Reviews           0
      Size              0
      Installs          0
      Type              0
      Price             0
      Content Rating    0
      dtype: int64
```

4.4 4.4. Changing Features Data Type

From the info of the features, we know that some features have ‘object’ Dtype with numeric value. So, we need to convert these features to ‘float’ dtype.

We noticed from the table shown that ‘Price’ column has ‘\$’ symbol, so this need to be dropped before converting ‘Price’ to float.

```
[14]: #Drop the $ sign from Price values
      df['Price'] = df['Price'].str.replace('$', '')
```

```
/var/folders/gy/jsxj5kgx6xj644yvj9136j480000gn/T/ipykernel_43877/2481516221.py:2
: FutureWarning: The default value of regex will change from True to False in a
future version. In addition, single character regular expressions will *not* be
treated as literal strings when regex=True.
      df['Price'] = df['Price'].str.replace('$', '')
```

```
[15]: #Converting Price to float
      df['Price'] = df['Price'].astype(float)
```

We also noticed from the table shown that “Installs” column has ‘+’ and ‘,’ symbols, so these need to be dropped before converting ‘Installs’ to float.

```
[16]: #Drop the + and , sign from Installs values
      df['Installs'] = df['Installs'].str.replace('+', '')
      df['Installs'] = df['Installs'].str.replace(',', '')
      df[['Installs']]
```

```
/var/folders/gy/jsxj5kgx6xj644yvj9136j480000gn/T/ipykernel_43877/283116142.py:2:
FutureWarning: The default value of regex will change from True to False in a
future version. In addition, single character regular expressions will *not* be
treated as literal strings when regex=True.
      df['Installs'] = df['Installs'].str.replace('+', '')
```

```
[16]:      Installs
      0      10000
```

```

1      500000
2      5000000
3      50000000
4      100000
...
9655      5000
9656      100
9657      1000
9658      1000
9659 10000000

```

[9658 rows x 1 columns]

```
[17]: #Converting Installs to float
df['Installs'] = df['Installs'].astype(float)
```

```
[18]: #Converting Reviews to float
df['Reviews'] = df['Reviews'].astype(float)
```

We also noticed from the table shown that “Size” column has ‘Varies with device’. This will be replaced with NaN in order to fill it with mean value later

```
[19]: df['Size'] = df['Size'].replace('Varies with device', 'NaN', regex=True)
```

Moreover, some apps have a size in K (kilobyte) and some have it in M (megabyte). We will convert the sizes to KB and MB. And we will convert the type to float

```
[20]: df['Size'] = (df['Size'].replace(r'[kM]+$', '', regex=True).astype(float) *
                  df['Size'].str.extract(r'[\d\.]+([kM]+)', expand=False)
                  .fillna(1).replace(['k','M'], [10**3, 10**6]).astype(int))
```

```
[21]: #Filling missing values with mean
SizegMean = df['Size'].mean()
df['Size'] = df['Size'].fillna(SizegMean)
```


5 5. Visualalization

We will take a glance into tha dataset in visuals.

```
[22]: #import visualization libraries
import seaborn as sns
import matplotlib.pyplot as plt
```

Lest’s first create a pairplot for the remainig numeric features of the dataset.


```
[23]: sns.pairplot(df);
```

Now we want to see the distribution of each column and the count of its categories.


5.0.1 Category

```
[24]: plt.rcParams['figure.figsize'] = (20, 10)
sns.countplot(x='Category', data=df)
plt.xticks(rotation=70);
```




5.0.2 Rating

```
[25]: plt.rcParams['figure.figsize'] = (20, 10)
RA = df['Rating'].value_counts()
plt.hist(RA, 50)
plt.xlabel('Count')
plt.ylabel('Rating')
plt.xticks(rotation=70);
```




5.0.3 Reviews

```
[26]: plt.rcParams['figure.figsize'] = (20, 10)
      R = df['Reviews'].value_counts()
      plt.hist(R, 100)
      plt.xlabel('Reviews')
      plt.ylabel('count')
      plt.xticks(rotation=70);
```




5.0.4 Size

```
[27]: plt.rcParams['figure.figsize'] = (20, 10)
      S= df['Size'].value_counts()
      plt.hist(S, 50)
      plt.xlabel('count')
      plt.ylabel('Size')
      plt.xticks(rotation=70);
```




5.0.5 Installs

```
[28]: plt.rcParams['figure.figsize'] = (20, 10)
sns.countplot(x='Installs',data=df)
plt.xticks(rotation=70);
```




5.0.6 Type

```
[29]: plt.rcParams['figure.figsize'] = (20, 10)
df['Type'].value_counts().plot.pie()
plt.xticks(rotation=70);
```




5.0.7 Price

```
[30]: plt.rcParams['figure.figsize'] = (20, 10)
P = df.loc[(df.Type == 'Paid')]
P['Price'].value_counts().plot.bar()
plt.xlabel('Price')
plt.ylabel('count')
plt.xticks(rotation=70);
```



5.0.8 Content Rating

```
[31]: plt.rcParams['figure.figsize'] = (20, 10)
sns.countplot(x='Content Rating',data=df)
plt.xticks(rotation=70);
```



6 6. Feature Engineering

We want to add a new column called (App_Demand) engineered engineered from the 'Installs' and 'Rating' features. The column will have 7 categories ('very_high_demand', 'high_demand', 'on_demand', 'moderate_demand', 'low_demand', 'very_low_demand', 'no_demand') which represent the demand on a certain app based on its installs and rating.

6.1 6.1. Calculations Related to New Column

We will compute maximum, average and minimum 'Installs' values, and average and minimum 'Rating' values using max(), mean(), and min() methods respectively. The approach compares the 'Installs' number and 'Rating' of each app with Maximum, Average, and Minimum Installs values and assigns the appropriate category for each app. These methods will be applied to the 'Installs' and 'Rating' features. In addition, we defined two variables (lowRateIN) which denotes the average 'Installs' value between the minimum and average values. And (highRateIN) which denotes the average 'Installs' value between the maximum and average values.

```
[32]: avgIN = round(df['Installs'].mean(), 2)
      maxiIN = round(df['Installs'].max(), 2)
      miniIN = round(df['Installs'].min(), 2)

      print ("Mean Value:", avgIN)
      print ("Max Value:", maxiIN)
      print ("Min Value:", miniIN)

      lowRateIN = round((miniIN+avgIN)/2, 2)
      print ("Below Average App - Cutoff:", lowRateIN)
      highRateIN = round((avgIN+maxiIN)/2, 2)
      print ("Above Average App - Cutoff:", highRateIN)
```

```
Mean Value: 7778312.02
Max Value: 1000000000.0
Min Value: 0.0
Below Average App - Cutoff: 3889156.01
Above Average App - Cutoff: 503889156.01
```

```
[33]: avgR = round(df['Rating'].mean(), 2)
      miniR = round(df['Rating'].min(), 2)

      print ("Mean Value:", avgR)
      print ("Min Value:", miniR)
```

```
Mean Value: 4.17
Min Value: 1.0
```

6.2 6.2. Inserting the New Column

We want to classify apps into seven categories. First, the 'very_high_demand' apps are the app with the installs number between maximum and highRateIN values. Second, the 'high_demand'

apps are the app with the installs number between highRateIN and average value and rating value above average value. Third, the 'on_demand' apps are the app with the installs number between highRateIN and average values and rating value below average value. Fourth, the 'moderate_demand' apps are the app with the installs number between average and lowrRateIN values. Fifth, 'low_demand' apps are the app with the installs number between lowrRateIN and minimum values and rating value above average value. Sixth, 'very_low_demand' apps are the app with the installs number between lowrRateIN and minimum values and rating value below average value. The last category, 'no_demand' apps are the app with the installs number equal to minimum values. The result is stored in a new column called (App_Demand) that is added to the dataset.

```
[34]: conditions = [
    (df['Installs'] <= maxiIN) & (df['Installs'] >= highRateIN),
    (df['Installs'] < highRateIN) & (df['Installs'] >= avgIN) & (df['Rating'] >
    ↪ avgR),
    (df['Installs'] < highRateIN) & (df['Installs'] >= avgIN) & (df['Rating']
    ↪ <= avgR),
    (df['Installs'] < avgIN) & (df['Installs'] >= lowRateIN),
    (df['Installs'] < lowRateIN) & (df['Installs'] > miniIN) & (df['Rating'] >
    ↪ avgR),
    (df['Installs'] < lowRateIN) & (df['Installs'] > miniIN) & (df['Rating'] <=
    ↪ avgR),
    (df['Installs'] == miniIN)]
categories = ['very_high_demand', 'high_demand', 'on_demand',
    ↪ 'moderate_demand', 'low_demand', 'very_low_demand',
    ↪ 'no_demand']
df['App_Demand'] = np.select(conditions, categories)

#lets show the new column on the first 15 rows
df.head(15)
```

```
[34]:
```

	Category	Rating	Reviews	Size	Installs	Type	Price	\
0	ART_AND_DESIGN	4.1	159.0	19000000.0	10000.0	Free	0.0	
1	ART_AND_DESIGN	3.9	967.0	14000000.0	500000.0	Free	0.0	
2	ART_AND_DESIGN	4.7	87510.0	8700000.0	5000000.0	Free	0.0	
3	ART_AND_DESIGN	4.5	215644.0	25000000.0	50000000.0	Free	0.0	
4	ART_AND_DESIGN	4.3	967.0	2800000.0	100000.0	Free	0.0	
5	ART_AND_DESIGN	4.4	167.0	5600000.0	50000.0	Free	0.0	
6	ART_AND_DESIGN	3.8	178.0	19000000.0	50000.0	Free	0.0	
7	ART_AND_DESIGN	4.1	36815.0	29000000.0	1000000.0	Free	0.0	
8	ART_AND_DESIGN	4.4	13791.0	33000000.0	1000000.0	Free	0.0	
9	ART_AND_DESIGN	4.7	121.0	3100000.0	10000.0	Free	0.0	
10	ART_AND_DESIGN	4.4	13880.0	28000000.0	1000000.0	Free	0.0	
11	ART_AND_DESIGN	4.4	8788.0	12000000.0	1000000.0	Free	0.0	
12	ART_AND_DESIGN	4.2	44829.0	20000000.0	10000000.0	Free	0.0	
13	ART_AND_DESIGN	4.6	4326.0	21000000.0	100000.0	Free	0.0	
14	ART_AND_DESIGN	4.4	1518.0	37000000.0	100000.0	Free	0.0	

	Content Rating	App_Demand
0	Everyone	very_low_demand
1	Everyone	very_low_demand
2	Everyone	moderate_demand
3	Teen	high_demand
4	Everyone	low_demand
5	Everyone	low_demand
6	Everyone	very_low_demand
7	Everyone	very_low_demand
8	Everyone	low_demand
9	Everyone	low_demand
10	Everyone	low_demand
11	Everyone	low_demand
12	Teen	high_demand
13	Everyone	low_demand
14	Everyone	low_demand

As we have seen, the new 'App_Demand' column is added to our dataset.

```
[35]: #Use loc function to test the results
df.loc[(df['Installs'] <= maxiIN) & (df['Installs'] >= highRateIN)].head(3)
```

```
[35]:
```

	Category	Rating	Reviews	Size	Installs \
152	BOOKS_AND_REFERENCE	3.9	1433233.0	2.039529e+07	1.000000e+09
299	COMMUNICATION	4.0	56642847.0	2.039529e+07	1.000000e+09
300	COMMUNICATION	4.4	69119316.0	2.039529e+07	1.000000e+09

	Type	Price	Content Rating	App_Demand
152	Free	0.0	Teen	very_high_demand
299	Free	0.0	Everyone	very_high_demand
300	Free	0.0	Everyone	very_high_demand

6.3. Visualization of the Distribution

Now we will display the distribution of the apps as per their categories of demand.

```
[36]: plt.figure(figsize=(10,6))
sns.set(context='notebook', style='whitegrid')
df['Index'] = df.index #converting index into column
sns.boxplot(x= 'Index', y= 'App_Demand', data= df, palette= 'flag')
sns.boxplot()
sns.swarmplot(x = 'Index', y= 'App_Demand', data= df, linewidth= 0)
```

```
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/seaborn/categorical.py:1296: UserWarning: 72.3% of the points cannot be
placed; you may want to decrease the size of the markers or use stripplot.
```

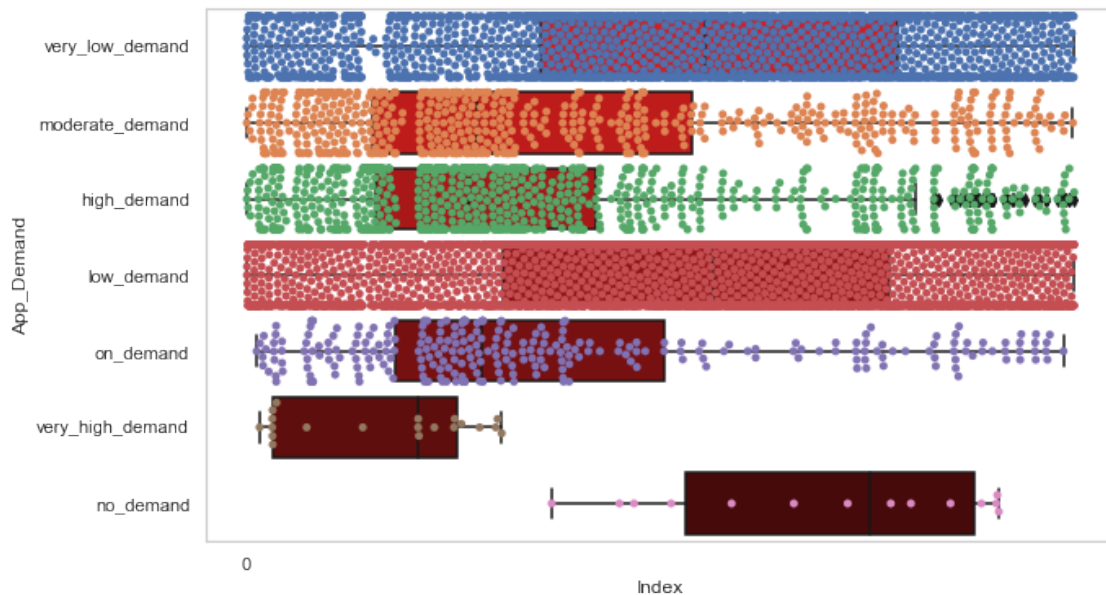
```
warnings.warn(msg, UserWarning)
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
```

```

packages/seaborn/categorical.py:1296: UserWarning: 32.3% of the points cannot be
placed; you may want to decrease the size of the markers or use stripplot.
  warnings.warn(msg, UserWarning)
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/seaborn/categorical.py:1296: UserWarning: 58.5% of the points cannot be
placed; you may want to decrease the size of the markers or use stripplot.
  warnings.warn(msg, UserWarning)
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/seaborn/categorical.py:1296: UserWarning: 85.0% of the points cannot be
placed; you may want to decrease the size of the markers or use stripplot.
  warnings.warn(msg, UserWarning)
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/seaborn/categorical.py:1296: UserWarning: 15.0% of the points cannot be
placed; you may want to decrease the size of the markers or use stripplot.
  warnings.warn(msg, UserWarning)

```

```
[36]: <AxesSubplot:xlabel='Index', ylabel='App_Demand'>
```



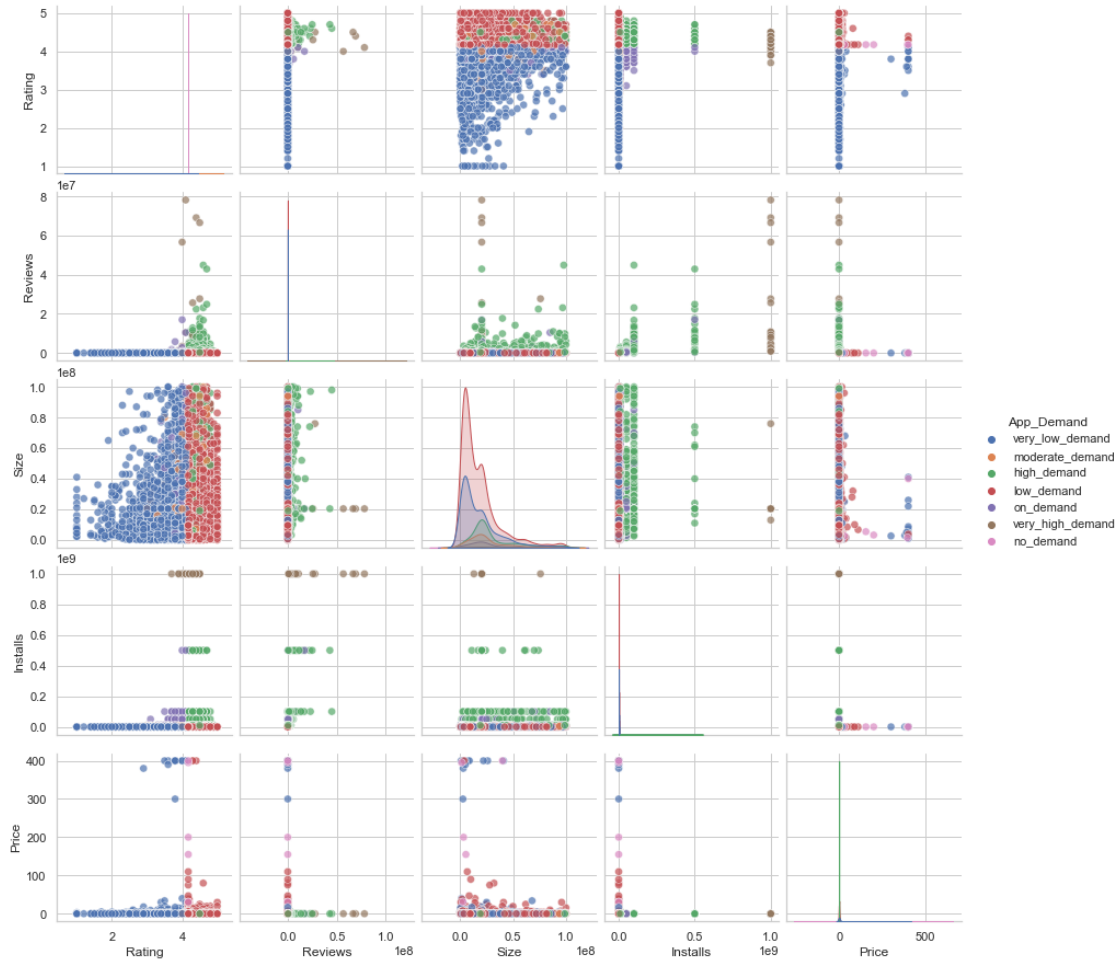
We will drop the 'Index' as we no longer need it.

```
[37]: df = df.drop(['Index'], axis = 1)
```

Lets see a pair plot of the dataset with the target App_Demand

```
[38]: sns.pairplot(data = df, hue = 'App_Demand', plot_kws={'alpha':0.7 , 's':50})
```

```
[38]: <seaborn.axisgrid.PairGrid at 0x7fb7d1165f40>
```



7. Data Transformation

We know that we have some features with string values. Let's take a look into those distinguish values

```
[39]: print(df['Category'].value_counts())
```

FAMILY	1831
GAME	959
TOOLS	827
BUSINESS	420
MEDICAL	395
PERSONALIZATION	376
PRODUCTIVITY	374
LIFESTYLE	369
FINANCE	345
SPORTS	325

COMMUNICATION	315
HEALTH_AND_FITNESS	288
PHOTOGRAPHY	281
NEWS_AND_MAGAZINES	254
SOCIAL	239
BOOKS_AND_REFERENCE	222
TRAVEL_AND_LOCAL	219
SHOPPING	202
DATING	171
VIDEO_PLAYERS	163
MAPS_AND_NAVIGATION	131
EDUCATION	119
FOOD_AND_DRINK	112
ENTERTAINMENT	102
AUTO_AND_VEHICLES	85
LIBRARIES_AND_DEMO	84
WEATHER	79
HOUSE_AND_HOME	74
EVENTS	64
ART_AND_DESIGN	64
PARENTING	60
COMICS	56
BEAUTY	53

Name: Category, dtype: int64

```
[40]: print(df['Type'].value_counts())
```

Free	8902
Paid	756

Name: Type, dtype: int64

```
[41]: print(df['Content Rating'].value_counts())
```

Everyone	7903
Teen	1036
Mature 17+	393
Everyone 10+	321
Adults only 18+	3
Unrated	2

Name: Content Rating, dtype: int64

```
[42]: print(df['App_Demand'].value_counts())
```

low_demand	5103
very_low_demand	2563
high_demand	1071
moderate_demand	607
on_demand	280

```

very_high_demand      20
no_demand              14
Name: App_Demand, dtype: int64

```

It is now important to encode categorical labels into numerical values. So, the values/labels in the 'Category', 'Type', 'Content Rating' and 'App_Demand' need to be transformed and normalized, such that they contain only numerical values. For instance, 'App_Demand' will contain only values between 0 and 6. For this purpose, we use the LabelEncoder from sklearn.preprocessing library and fit the categories on it to get the numeric values.

```

[43]: from sklearn.preprocessing import LabelEncoder
      encode = LabelEncoder()

      df['Category'] = encode.fit_transform(df['Category'])

      df['Type'] = encode.fit_transform(df['Type'])

      df['Content Rating'] = encode.fit_transform(df['Content Rating'])

      df['App_Demand'] = encode.fit_transform(df['App_Demand'])

```

```

[44]: df.head(3)

```

```

[44]:
   Category  Rating  Reviews      Size  Installs  Type  Price  \
0         0    4.1    159.0  19000000.0   10000.0    0    0.0
1         0    3.9    967.0  14000000.0   500000.0    0    0.0
2         0    4.7  87510.0   8700000.0  5000000.0    0    0.0

   Content Rating  App_Demand
0              1           6
1              1           6
2              1           2


```

8 8. Correlation Matrix

```

[45]: plt.figure(figsize=(10,10))
      cormat = df.corr()
      sns.heatmap(cormat, annot=True, cmap='mako', square=True);

```



We can notice that almost all the features have no relationship with the target of classification ‘App_Demand’. We can already assume that the classification models won’t work. But we will continue and see the results.

We can also see that the target of regression model ‘Rating’ has low relationship with the other features. We can also assume that we will have bad regression.

9 9. Building the Models

First, we plan to develop K-Nearest Neighbor and Random Forrest Classifier models for classifying the apps into categories based on their demand. Second, we plan to develop a Random Forrest Regression model to predict the app rating.

9.1 9.1. Imorting the Libraries

```
[46]: from sklearn.model_selection import train_test_split
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.ensemble import RandomForestRegressor
```

```
[47]: from sklearn import metrics
      from sklearn.metrics import accuracy_score
      from sklearn.metrics import classification_report
      from sklearn.metrics import mean_squared_error
      from sklearn.metrics import confusion_matrix
```

9.2 9.2. Classification Goal

9.2.1 9.2.1. Set Features and Target

We will specify X set of features and y feature which is the target. We will use features ('Category', 'Type', 'Price', 'Content Rating') as input variables, and the new feature (App_Demand) will be generated as output for our classification models. X will take all rows, and ('Category', 'Type', 'Price', 'Content Rating') columns. And y will take all rows, and only the last column 'App_Demand'.

```
[48]: X = df[['Category', 'Type', 'Price', 'Content Rating']]
      y = df.iloc[:, -1]
```

```
[49]: X.head()
```

```
[49]:
```

	Category	Type	Price	Content Rating
0	0	0	0.0	1
1	0	0	0.0	1
2	0	0	0.0	1
3	0	0	0.0	4
4	0	0	0.0	1

```
[50]: y.head()
```

```
[50]:
```

0	6
1	6
2	2
3	0
4	1

Name: App_Demand, dtype: int64

9.2.2 9.2.2. Splitting the Dataset

The dataset must be split into a training set and testing set. We will use the `train_test_split()` method to do splitting. We will set the `test_size` parameter 0.30.

```
[51]: #hold out 30% of the data for final testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.30,
↳random_state=42, stratify=y)
```

9.2.3 9.2.3. Feature Scaling

Before fitting the model, we need to scales and translates each feature on the training set to be in range between 0 and 1.

```
[52]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
```

```
[53]: scaler.fit(X_train)
```

```
[53]: MinMaxScaler()
```

9.2.4 9.2.4. Training and Evaluating the Models

9.2.5 9.2.4.1. K-Nearest Neighbor Classifier

Best Value of K Tuning Using Cross-Validation

```
[54]: #search for an optimal value of K for KNN
from sklearn.model_selection import cross_val_score

neighbors = list(range(0, 300, 5))
k_scores = []
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')
    k_scores.append(scores.mean())
print(k_scores)
```

/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator fit failed. The score on this train-test partition for these parameters will be set to nan. Details:

Traceback (most recent call last):

File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py", line 179, in fit
return self._fit(X, y)

File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_base.py", line 514, in _fit
raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0


```
warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
```

```
ValueError: Expected n_neighbors > 0. Got 0
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
```

```
ValueError: Expected n_neighbors > 0. Got 0
```

```
warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
```

```
Traceback (most recent call last):
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
```

```
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
```

```

    raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0

    warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0

```

```

    warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0

```

```

    warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit

```

```

    return self._fit(X, y)
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0

warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0

warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0

warnings.warn("Estimator fit failed. The score on this train-test"
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py:615: FitFailedWarning: Estimator
fit failed. The score on this train-test partition for these parameters will be
set to nan. Details:
Traceback (most recent call last):
  File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/model_selection/_validation.py", line 598, in _fit_and_score

```

```

    estimator.fit(X_train, y_train, **fit_params)
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_classification.py", line 179, in fit
    return self._fit(X, y)
File "/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/neighbors/_base.py", line 514, in _fit
    raise ValueError(
ValueError: Expected n_neighbors > 0. Got 0

```

```
warnings.warn("Estimator fit failed. The score on this train-test"
```

```

[nan, 0.4625739644970414, 0.4936390532544378, 0.5078402366863906,
0.5131656804733727, 0.5178994082840236, 0.5170118343195267, 0.5173076923076924,
0.5202662721893491, 0.5217455621301774, 0.5239644970414201, 0.5264792899408283,
0.525887573964497, 0.5235207100591716, 0.5235207100591717, 0.5230769230769232,
0.5238165680473373, 0.5261834319526627, 0.5260355029585798, 0.5247041420118344,
0.5235207100591716, 0.5247041420118344, 0.5245562130177515, 0.5254437869822484,
0.5267751479289939, 0.5273668639053255, 0.5275147928994082, 0.528698224852071,
0.5269230769230768, 0.5295857988165681, 0.5309171597633137, 0.5292899408284024,
0.5284023668639053, 0.5288461538461539, 0.5279585798816567, 0.5266272189349113,
0.5267751479289939, 0.526775147928994, 0.5272189349112426, 0.5267751479289942,
0.5285502958579882, 0.5281065088757397, 0.5278106508875741, 0.5270710059171598,
0.5272189349112427, 0.5270710059171598, 0.5270710059171598, 0.526923076923077,
0.5273668639053255, 0.5275147928994084, 0.5273668639053255, 0.5275147928994084,
0.5276627218934913, 0.5275147928994084, 0.5275147928994084, 0.5278106508875741,
0.527958579881657, 0.5282544378698226, 0.5284023668639054, 0.5284023668639054]


```

```

[55]: # plot the value of K for KNN (x-axis) versus the cross-validated accuracy
      ↪ (y-axis)
import matplotlib.pyplot as plt
plt.plot(neighbors, k_scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')

```

```
[55]: Text(0, 0.5, 'Cross-Validated Accuracy')
```



The best result is nearly at $k=145$

Building the Model (1st Trial)

```
[56]: knn1 = KNeighborsClassifier(n_neighbors=150)
      knn1.fit(X_train, y_train)
```

```
[56]: KNeighborsClassifier(n_neighbors=150)
```

```
[57]: y_predknn1 = knn1.predict(X_test)
```

```
[58]: knn1.predict([[4, 1, 3, 2]])
```

```
[58]: array([1])
```

Evaluating the 1st Trial

```
[59]: print("KNN Training Set Accuracy:", metrics.accuracy_score(y_train, knn1.
      ↪predict(X_train)))
      print("KNN Testing Set Accuracy:", metrics.accuracy_score(y_test, y_predknn1))
      print("Classification Report of KNN:\n", classification_report(y_test,
      ↪y_predknn1))
      print("Confusion Matrix of KNN:\n", confusion_matrix(y_test, y_predknn1))
```

KNN Training Set Accuracy: 0.531508875739645

KNN Testing Set Accuracy: 0.525879917184265

Classification Report of KNN:

	precision	recall	f1-score	support
0	0.27	0.10	0.15	322

1	0.54	0.97	0.69	1531
2	0.00	0.00	0.00	182
3	0.00	0.00	0.00	4
4	0.00	0.00	0.00	84
5	0.00	0.00	0.00	6
6	0.00	0.00	0.00	769
accuracy			0.53	2898
macro avg	0.12	0.15	0.12	2898
weighted avg	0.31	0.53	0.38	2898

Confusion Matrix of KNN:

```
[[ 33 289   0   0   0   0   0]
 [ 38 1491   0   0   0   0   2]
 [   9  173   0   0   0   0   0]
 [   0   4   0   0   0   0   0]
 [   7   77   0   0   0   0   0]
 [   1    5   0   0   0   0   0]
 [  35  734   0   0   0   0   0]]
```

```
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

As shown in the confusion matrix; there is clearly a bias towards the class of 1. Which means that the dataset is maybe imbalanced.

Rsampling the Dataset Using SMOTE - Synthetic Minority Oversampling TEchnique

Let's take a look to the target class of our dataset

```
[60]: y.value_counts()
```

```
[60]: 1    5103
      6    2563
      0    1071
      2     607
```

```
4      280
5      20
3      14
Name: App_Demand, dtype: int64
```

```
[61]: y_train.value_counts()
```

```
[61]: 1      3572
      6      1794
      0      749
      2      425
      4      196
      5       14
      3       10
      Name: App_Demand, dtype: int64
```

We can notice that clearly the dataset is imbalanced. We will apply SMOTE - Synthetic Minority Oversampling TEchnique on the imbalanced dataset to resample it.

```
[62]: from imblearn.over_sampling import SMOTE
      sm=SMOTE(random_state=42)
      X_smoted, y_smoted =sm.fit_resample(X_train,y_train)
```

Now let's check the target class value counts

```
[63]: y_smoted.value_counts()
```

```
[63]: 6      3572
      2      3572
      1      3572
      4      3572
      0      3572
      5      3572
      3      3572
      Name: App_Demand, dtype: int64
```

Best Value of K Tuning Using Cross-Validation


```
[64]: #search for an optimal value of K for KNN
      neighbors = list(range(1, 1000, 10))
      k_scores = []
      for k in neighbors:
          knn = KNeighborsClassifier(n_neighbors=k)
          scores = cross_val_score(knn, X_smoted, y_smoted, cv=10, scoring='accuracy')
          k_scores.append(scores.mean())
      print(k_scores)
```

```
[0.3382266613354658, 0.360103606557377, 0.37246163934426224,
```

0.37474159136345464, 0.3799400719712115, 0.3759011275489804,
0.37382175129948025, 0.3727418632546981, 0.3705019432227109, 0.373061287485006,
0.37198117552978804, 0.3740205357856857, 0.3746204398240704,
0.37690026389444226, 0.37486040783686525, 0.3735405517792883,
0.3777398320671731, 0.3808987604958016, 0.3814993522590963, 0.3836189044382247,
0.383058856457417, 0.3840186005597761, 0.3832587445021991, 0.3817790323870452,
0.380259112355058, 0.37897903238704517, 0.37677964014394244, 0.3757395761695322,
0.3730601999200319, 0.3685009036385446, 0.3657812235105958, 0.36450136745301875,
0.3630215753698521, 0.35782261495401835, 0.3552629348260696, 0.3545431267493003,
0.3545830787684926, 0.35374311075569775, 0.3551427588964414, 0.3549828868452619,
0.35470299880047984, 0.35610287085165926, 0.35662267892842864,
0.35750240703718517, 0.35822253498600554, 0.3596225349860056,
0.36110266293482607, 0.3598229028388644, 0.35938293482606953,
0.3594229348260696, 0.35926295081967213, 0.3597426949220312, 0.360062662934826,
0.3593828068772491, 0.35946271091563375, 0.35826266293482606,
0.35766312674930034, 0.35474322271091563, 0.35430347860855654,
0.3530235425829668, 0.3546232227109156, 0.35498315873650543,
0.35426341463414635, 0.35794251899240304, 0.367460343862455, 0.3691807277089164,
0.3703809676129548, 0.37002087165133946, 0.3694609676129548, 0.3698609356257497,
0.36982083966413437, 0.37258042383046786, 0.37226058376649335,
0.37246067972810876, 0.3719010475809676, 0.3727005517792883, 0.3731403758496601,
0.37234048780487805, 0.37234063174730114, 0.3710608716513395,
0.3704610155937625, 0.37138061575369846, 0.3712605677728909, 0.3703004238304678,
0.37066034386245506, 0.3698202798880447, 0.36842042383046775,
0.36754035985605754, 0.36702029588164736, 0.36318167133146745,
0.3624216393442623, 0.36118192722910836, 0.35994187924830073,
0.35898203918432625, 0.35910223110755696, 0.3583021031587365,
0.35630259896041583, 0.3541028388644542, 0.3521029668132747,
0.35186296681327467]

```
[65]: # plot the value of K for KNN (x-axis) versus the cross-validated accuracy
      ↪ (y-axis)
import matplotlib.pyplot as plt
plt.plot(neighbors, k_scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')
```

```
[65]: Text(0, 0.5, 'Cross-Validated Accuracy')
```

Building the Model (2nd Trial) We will build the model now using resampled dataset

```
[66]: knn2 = KNeighborsClassifier(n_neighbors=220)
      knn2.fit(X_smoted, y_smoted)
```

```
[66]: KNeighborsClassifier(n_neighbors=220)
```

```
[67]: y_predknn2 = knn2.predict(X_test)
```

Evaluating the 2nd Trial

```
[68]: print("KNN Training Set Accuracy:", metrics.accuracy_score(y_smoted, knn2.
      ↪predict(X_smoted)))
      print("KNN Testing Set Accuracy:", metrics.accuracy_score(y_test, y_predknn2))
      print("Classification Report of KNN:\n", classification_report(y_test,
      ↪y_predknn2))
      print("Confusion Matrix of KNN:\n", confusion_matrix(y_test, y_predknn2))
```

KNN Training Set Accuracy: 0.3874980003199488

KNN Testing Set Accuracy: 0.3053830227743271

Classification Report of KNN:

	precision	recall	f1-score	support
0	0.19	0.61	0.28	322
1	0.61	0.41	0.49	1531
2	0.09	0.17	0.12	182
3	0.01	0.75	0.03	4
4	0.04	0.02	0.03	84
5	0.01	0.17	0.02	6

	6	0.32	0.04	0.07	769
accuracy				0.31	2898
macro avg	0.18	0.31	0.15		2898
weighted avg	0.43	0.31	0.32		2898

Confusion Matrix of KNN:

```
[[196  49  44   1   9  20   3]
 [445 623 176 170  16  47  54]
 [ 78  52  31   0   7  11   3]
 [  0   0   0   3   0   0   1]
 [ 47  19   9   0   2   7   0]
 [  4   1   0   0   0   1   0]
 [287 282  95  42  11  23  29]]
```

We somehow fixed the imbalance of the dataset, but the the accuracy is worse.

9.2.6 9.2.4.2. Random Forrest Classifier

Building the Model (1st Trial)

```
[69]: RF1 = RandomForestClassifier(n_estimators = 1000, criterion = 'entropy',
    ↪random_state = 42)
RF1.fit(X_train, y_train)
```

```
[69]: RandomForestClassifier(criterion='entropy', n_estimators=1000, random_state=42)
```

```
[71]: y_predRF1 = RF1.predict(X_test)
```

Evaluating the 1st Trial

```
[72]: print("RF Training Set Accuracy:", metrics.accuracy_score(y_train, RF1.
    ↪predict(X_train)))
print("RF Testing Set Accuracy:", metrics.accuracy_score(y_test, y_predRF1))
print("Classification Report of RF:\n", classification_report(y_test,
    ↪y_predRF1))
print("Confusion Matrix of RF:\n", confusion_matrix(y_test, y_predRF1))
```

RF Training Set Accuracy: 0.5455621301775148

RF Testing Set Accuracy: 0.5158730158730159

Classification Report of RF:

	precision	recall	f1-score	support
0	0.25	0.12	0.16	322
1	0.54	0.92	0.68	1531
2	0.00	0.00	0.00	182
3	0.00	0.00	0.00	4
4	0.00	0.00	0.00	84
5	0.00	0.00	0.00	6

	6	0.31	0.06	0.10	769
accuracy				0.52	2898
macro avg		0.16	0.16	0.14	2898
weighted avg		0.40	0.52	0.41	2898

Confusion Matrix of RF:

```
[[ 39 266  1  0  0  0 16]
 [ 51 1411 1  0  0  0 68]
 [ 14 155  0  0  0  0 13]
 [  0  4  0  0  0  0  0]
 [ 13  66  0  0  0  0  5]
 [  2  4  0  0  0  0  0]
 [ 40 683 1  0  0  0 45]]
```

```
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Users/hessaalhamad/opt/anaconda3/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1248: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

The Random Forrest performed almost as accurate as KNN but with no bias towards the class of 1. But still, there are some missing classes. So, the model must be build on balanced dataset.

Building the Model (2nd Trial) We will use (class weight) in Random Forest classifier to see if it solve the problem of imbalanced dataset

```
[73]: RF2 = RandomForestClassifier(n_estimators = 1000, criterion = 'entropy',
    ↪random_state = 42, class_weight= 'balanced')
RF2.fit(X_train, y_train)
```

```
[73]: RandomForestClassifier(class_weight='balanced', criterion='entropy',
    n_estimators=1000, random_state=42)
```

```
[74]: y_predRF2 = RF2.predict(X_test)
```

Evaluating the 2nd Trial

```
[75]: print("RF Training Set Accuracy:", metrics.accuracy_score(y_train, RF2.
        ↪predict(X_train)))
print("RF Testing Set Accuracy:", metrics.accuracy_score(y_test, y_predRF2))
print("Classification Report of RF:\n", classification_report(y_test,
        ↪y_predRF2))
print("Confusion Matrix of RF:\n", confusion_matrix(y_test, y_predRF2))
```

RF Training Set Accuracy: 0.28816568047337277

RF Testing Set Accuracy: 0.25086266390614215

Classification Report of RF:

	precision	recall	f1-score	support
0	0.25	0.43	0.32	322
1	0.73	0.21	0.33	1531
2	0.07	0.13	0.09	182
3	0.00	0.00	0.00	4
4	0.05	0.24	0.08	84
5	0.02	0.50	0.03	6
6	0.29	0.29	0.29	769
accuracy			0.25	2898
macro avg	0.20	0.26	0.16	2898
weighted avg	0.50	0.25	0.29	2898

Confusion Matrix of RF:

```
[[139  8  51  4  51  30  39]
 [217 321 179 125 191  62 436]
 [ 40  9  24  4  32  27  46]
 [  0  3  0  0  0  0  1]
 [ 25  3  9  1  20  10  16]
 [  2  0  0  0  1  3  0]
 [129 96  80  52 138  54 220]]
```

We somehow fixed the imbalance of the dataset, but the accuracy is also worse.

Building the Model (3rd Trial) We will build the model now using resampled dataset

```
[76]: RF3 = RandomForestClassifier(n_estimators = 1000, criterion = 'entropy',
        ↪random_state = 42)
RF3.fit(X_smoted, y_smoted)
```

```
[76]: RandomForestClassifier(criterion='entropy', n_estimators=1000, random_state=42)
```

```
[77]: y_predRF3 = RF3.predict(X_test)
```

Evaluating the 3rd Trial

```
[78]: print("RF Training Set Accuracy:", metrics.accuracy_score(y_smoted, RF2.
        ↪predict(X_smoted)))
print("RF Testing Set Accuracy:", metrics.accuracy_score(y_test, y_predRF3))
print("Classification Report of RF:\n", classification_report(y_test,
        ↪y_predRF3))
print("Confusion Matrix of RF:\n", confusion_matrix(y_test, y_predRF3))
```

RF Training Set Accuracy: 0.28943369060950247

RF Testing Set Accuracy: 0.28778467908902694

Classification Report of RF:

	precision	recall	f1-score	support
0	0.25	0.42	0.31	322
1	0.71	0.28	0.40	1531
2	0.07	0.13	0.09	182
3	0.11	0.75	0.19	4
4	0.06	0.29	0.09	84
5	0.02	0.50	0.03	6
6	0.29	0.28	0.29	769
accuracy			0.29	2898
macro avg	0.21	0.38	0.20	2898
weighted avg	0.49	0.29	0.33	2898

Confusion Matrix of RF:

```
[[135 12 53 2 51 34 35]
 [219 429 181 18 183 72 429]
 [ 43 14 24 0 34 22 45]
 [ 0 1 0 3 0 0 0]
 [ 23 3 9 0 24 9 16]
 [ 1 0 1 0 1 3 0]
 [129 141 79 4 143 57 216]]
```

Although the testing set performed better than the second trial; the results are still not good.

9.3 Regression Goal

Data Transformation Before building the regression model, we need to log-transform the ‘Installs’ column, in order to make it more ‘Normal’. Since there are apps with 0 installs, we will transform it with $\log(x+1)$ transform.

```
[79]: df['Installs']=np.log(df['Installs'] + 1)
```

Target Values Now, let’s check the value counts of the ‘Rating’ column

```
[80]: df['Rating'].value_counts()
```

```

[80]: 4.175052    1462
      4.300000     897
      4.400000     895
      4.500000     848
      4.200000     810
      4.600000     683
      4.100000     621
      4.000000     513
      4.700000     442
      3.900000     359
      3.800000     286
      5.000000     271
      3.700000     224
      4.800000     221
      3.600000     167
      3.500000     156
      3.400000     126
      3.300000     100
      4.900000      85
      3.000000      81
      3.100000      69
      3.200000      63
      2.900000      45
      2.800000      40
      2.600000      24
      2.700000      23
      2.500000      20
      2.300000      20
      2.400000      19
      1.000000      16
      2.200000      14
      2.000000      12
      1.900000      11
      1.700000       8
      1.800000       8
      2.100000       8
      1.600000       4
      1.400000       3
      1.500000       3
      1.200000       1
      Name: Rating, dtype: int64

```

There is only one data point with the value '1.200000'. This could make the learning model bias and not fit. So, we need to drop it.

```

[81]: #We need to locate it
      df.loc[(df['Rating'] == 1.200000)]

```

```
[81]:
```

	Category	Rating	Reviews	Size	Installs	Type	Price	\
8922	12	1.2	44.0	27000000.0	6.908755	0	0.0	

	Content Rating	App_Demand
8922	1	6

```
[82]: #Need to be dropped
df= df.drop(8922)
```

```
[83]: #Check the value counts again
df['Rating'].value_counts()
```

```
[83]: 4.175052    1462
      4.300000     897
      4.400000     895
      4.500000     848
      4.200000     810
      4.600000     683
      4.100000     621
      4.000000     513
      4.700000     442
      3.900000     359
      3.800000     286
      5.000000     271
      3.700000     224
      4.800000     221
      3.600000     167
      3.500000     156
      3.400000     126
      3.300000     100
      4.900000      85
      3.000000      81
      3.100000      69
      3.200000      63
      2.900000      45
      2.800000      40
      2.600000      24
      2.700000      23
      2.500000      20
      2.300000      20
      2.400000      19
      1.000000      16
      2.200000      14
      2.000000      12
      1.900000      11
      1.700000       8
      1.800000       8
```

```

2.100000      8
1.600000      4
1.400000      3
1.500000      3
Name: Rating, dtype: int64

```

9.3.1 9.3.1. Set Features and Target

We will specify X set of fetures and y feature which is the target. We will all remaining features exept ('Rating' and 'App_Demand') as input variables, and the feature (Rating) will be generated as output for our regression model. X will take all rows, and all columns expect tha new one 'App_Demand' and the 'Rating' column because it's the output(y). And y will take all rows, and only the column 'Rating'.

```

[84]: X1 =df.drop(['Rating','App_Demand'] , axis = 1)
      y1 =df['Rating']

```

9.3.2 9.3.2. Splitting the Dataset

The dataset must be split into a training set and testing set. We will use the train_test_split() method to do splitting. We will set the test_size parameter 0.30 for now, later on, we will run many experiments with different amounts of the test_size parameter.

```

[85]: #hold out 30% of the data for final testing
      X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, test_size=.30,
      ↪random_state=42)

```

9.3.3 9.3.3. Features Scaling

Before fitting the model, we need to scales and translates each feature on the training set to be in range between 0 and 1.

```

[86]: scaler = MinMaxScaler()

```

```

[87]: scaler.fit(X_train1)

```

```

[87]: MinMaxScaler()

```

```

[89]: X_train1 = scaler.transform(X_train1)
      X_test1 = scaler.transform(X_test1)

```

9.3.4 9.3.4. Training and Evaluating the Models

Building the Model

```

[90]: RFR1 = RandomForestRegressor()
      RFR1.fit(X_train1,y_train1)

```



```
[90]: RandomForestRegressor()
```

```
[91]: y_predRFR1 = RFR1.predict(X_test1)
```

```
[92]: RFR1.predict([[4, 3, 4, 50000, 1, 3, 2]])
```

```
[92]: array([4.427])
```

Evaluating the Model

```
[93]: mseRFR1 = mean_squared_error(y_test1, y_predRFR1)
      print("RMSE using RFR: ", np.sqrt(mseRFR1))
```

RMSE using RFR: 0.46971027534528054

9.3.5 9.3.5. Different Splitting of the Dataset.

In these experiments, we will change the dataset split size for the regression model each time to see if the results are affected by the way of split or not.

Training set > Testing set (15)

```
[96]: #hold out 15% of the data for final testing
      X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, test_size=.15,
      ↪random_state=42)
```

```
[97]: X_train1 = scaler.transform(X_train1)
      X_test1 = scaler.transform(X_test1)
```

```
[98]: RFR2 = RandomForestRegressor()
      RFR2.fit(X_train1, y_train1)
```

```
[98]: RandomForestRegressor()
```

```
[99]: y_predRFR2 = RFR2.predict(X_test1)
```

```
[100]: mseRFR2 = mean_squared_error(y_test1, y_predRFR2)
      print("RMSE using RFR: ", np.sqrt(mseRFR2))
```

RMSE using RFR: 0.45493198981234695

Training set = Testing set (50)

```
[101]: #hold out 50% of the data for final testing
      X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, test_size=.50,
      ↪random_state=42)
```

```
[102]: X_train1 = scaler.transform(X_train1)
      X_test1 = scaler.transform(X_test1)
```

```
[103]: RFR3 = RandomForestRegressor()  
RFR3.fit(X_train1,y_train1)
```

```
[103]: RandomForestRegressor()
```

```
[104]: y_predRFR3 = RFR3.predict(X_test1)
```

```
[105]: mseRFR3 = mean_squared_error(y_test1, y_predRFR3)  
print("RMSE using RFR: ", np.sqrt(mseRFR3))
```

RMSE using RFR: 0.49064546580455526

Training set < Testing set (75)

```
[106]: #hold out 75% of the data for final testing  
X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, test_size=.75,   
↳random_state=42)
```

```
[107]: X_train1 = scaler.transform(X_train1)  
X_test1 = scaler.transform(X_test1)
```

```
[108]: RFR4 = RandomForestRegressor()  
RFR4.fit(X_train1,y_train1)
```

```
[108]: RandomForestRegressor()
```

```
[109]: y_predRFR4 = RFR4.predict(X_test1)
```

```
[110]: mseRFR4 = mean_squared_error(y_test1, y_predRFR4)  
print("RMSE using RFR: ", np.sqrt(mseRFR4))
```

RMSE using RFR: 0.49109259386519327


Now let's visualize the testing set size VS. regression RMSE

```
[111]: # x axis values  
a = [25, 35, 50, 75] #the test set size  
#corresponding y axis values  
b = [np.sqrt(mseRFR2), np.sqrt(mseRFR1), np.sqrt(mseRFR3), np.sqrt(mseRFR4)]   
↳#refression results  
  
plt.grid(color='grey', linestyle='-', linewidth=0.25, alpha=0.5)  
  
plt.plot(a, b, color='red', linestyle='dashed', linewidth = 1,  
         marker='o', markerfacecolor='blue', markersize=8)  
  
#naming the x axis
```

```
plt.xlabel('The amount of test set')
# naming the y axis
plt.ylabel('Root mean squared error')

#giving a title to my graph
plt.title('impact of split set size on the regression results')

#function to show the plot
plt.show()
```



The above figure shows that the RMSE is at its lowest at the less test size which means that it gave the best regression.

10 10. Tools

In this project, we used the following tools: - Anaconda Navigator 2.1.1/ Jupyter Notebook 6.4.5 for implementing both algorithms and creating the models. - Set of libraries for modeling and visualization. In principle we are going to import these libraries : (Panda, Numpy, Matplotlib, Seaborn, Scikit-learn).

11 11. Conclusions

Depending on this dataset we could not predict the demand of apps based on their category, type, price and content rating. We attribute that to some reasons:

- The dataset is highly imbalanced.
- There were no strong relationships among features.

Regarding the predicting of the rating, we –somehow- reach some reasonable results. For further, we suggest building different regression algorithms to find the best result.